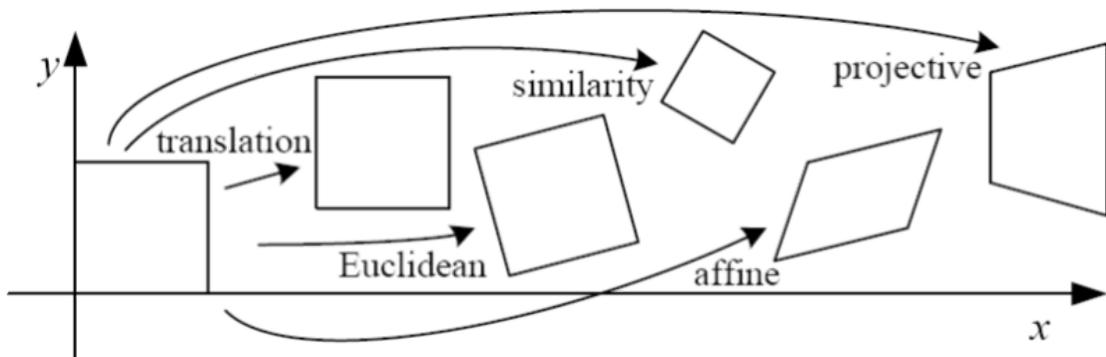




Tutorial 06 - Homography, Alignment & Panoramas



Agenda

- Matching Local Features
- Parametric Transformations
- Computing Parametric Transformations
 - Affine
 - Projective
 - Excercise 1 (Spring 2020 Q1A))
- RANSAC
 - Excercise 2 (Spring 2020 Q1D) - Ransac---Ransac)
- Panorama
 - Warping
 - Excercise 3 (Spring 2020 Q1B))
 - Image Blending (Feathering)
- Kornia & Transformations in Deep Learning
- Recommended Videos
- Credits

The largest panorama in the world (2014): Mont Blanc

In2WHITE Video

In2WHITE Full Image

Homographic usage examples:

<http://blog.flickr.net/en/2010/01/27/a-look-into-the-past/>

<https://www.instagram.com/albumplusart/>

In [2]:

```
%%html
<iframe src="http://www.in2white.com/" width="700" height="600"></iframe>
```



Filippo B.

▲ Sponsored by: ▲

©2015

In [1]:

```
# imports for the tutorial
import numpy as np,sys
import matplotlib.pyplot as plt
from PIL import Image
import cv2
from scipy import signal
%matplotlib inline
```

In [3]:

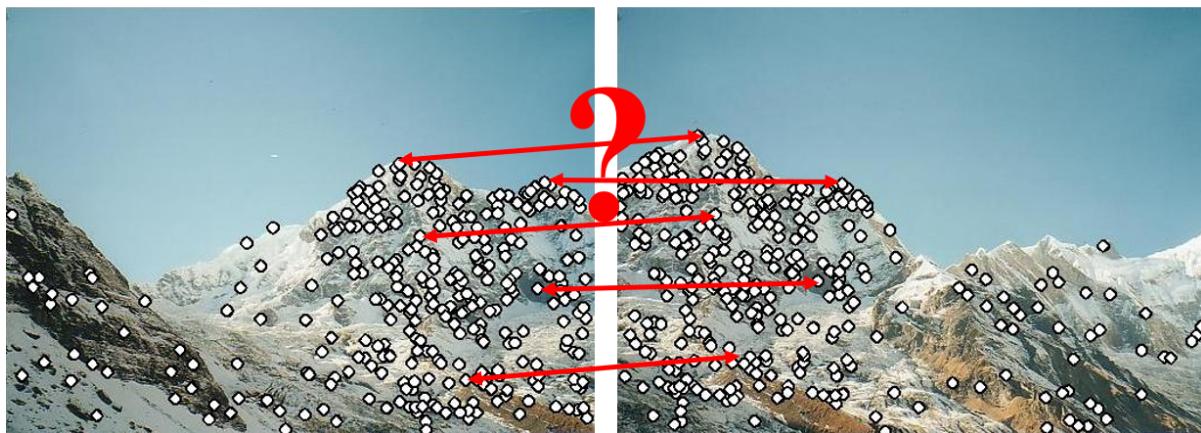
```
# plot images function
def plot_images(image_list, title_list, subplot_shape=(1,1), axis='off', fontsize=30, fi
    plt.figure(figsize=figsize)
    for ii, im in enumerate(image_list):
        c_title = title_list[ii]
        if len(cmap) > 1:
            c_cmap = cmap[ii]
        else:
            c_cmap = cmap[0]
        plt.subplot(subplot_shape[0], subplot_shape[1],ii+1)
        plt.imshow(im, cmap=c_cmap)
        plt.title(c_title, fontsize=fontsize)
        plt.axis(axis)
```



Matching Local Features

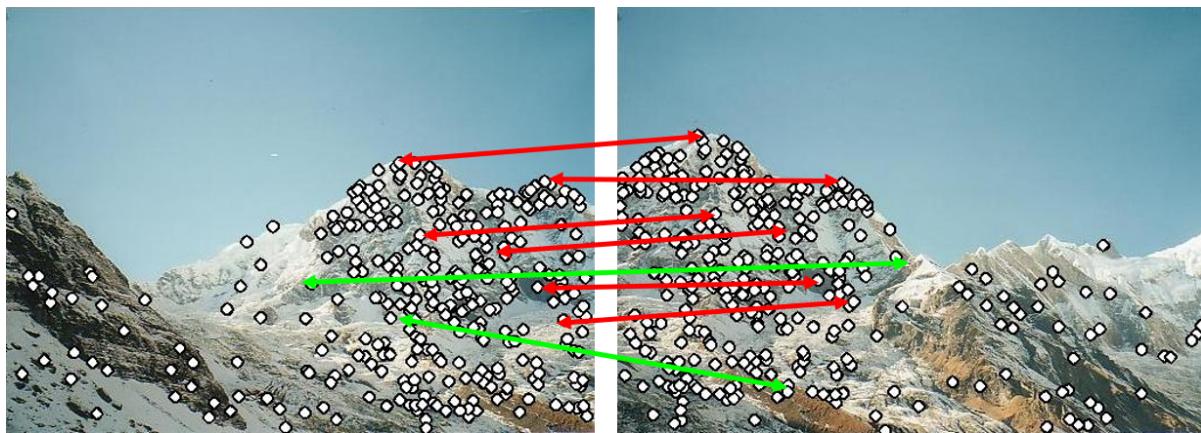
Feature matching

- We know how to detect and describe good points
- Next question: How to match them?



Typical feature matching results

- Some matches are correct
- Some matches are incorrect



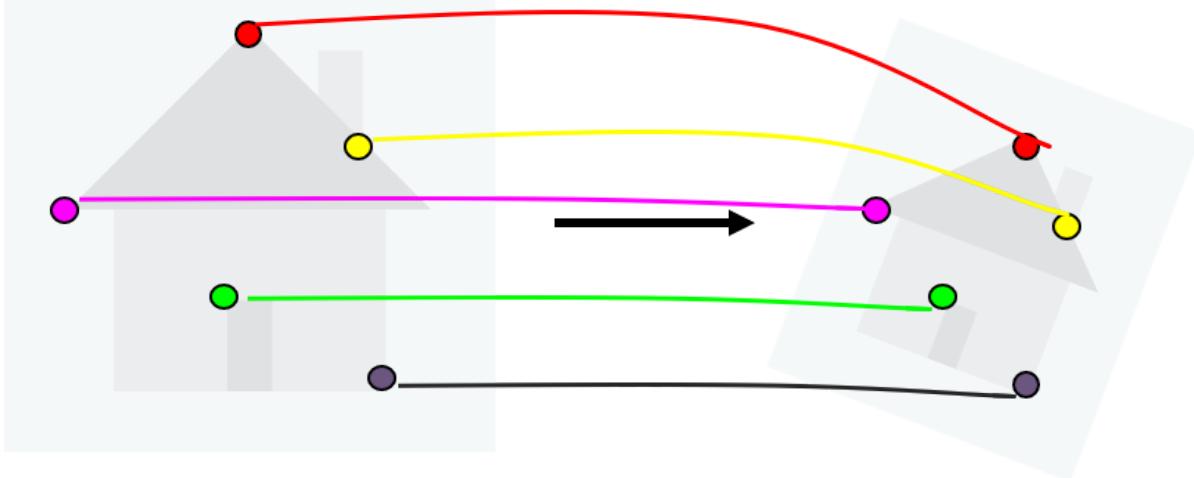
- Solution: search for a set of geometrically consistent matches



Parametric Transformations

Image Alignment

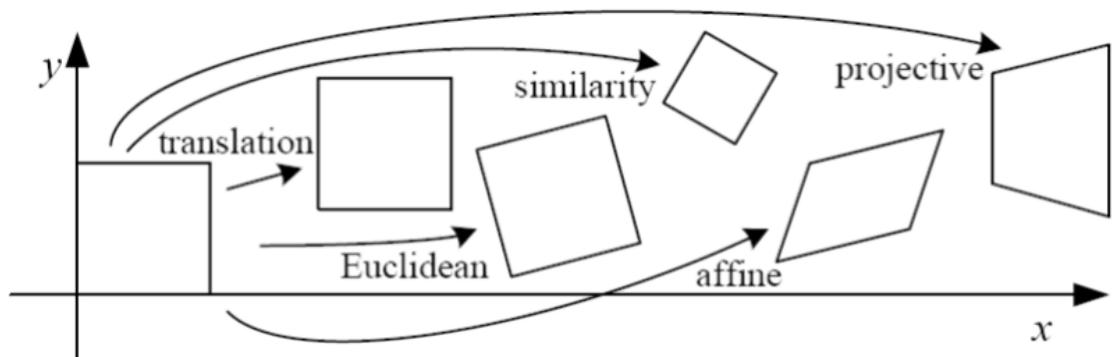
Given a set of matches, what parametric model describes a geometrically consistent transformation?



Basic 2D Transformations

Transformation	Matrix	# DoF	Preserves	Icon
translation	$\left[\begin{array}{c c} \mathbf{I} & \mathbf{t} \end{array} \right]_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\left[\begin{array}{c c} \mathbf{R} & \mathbf{t} \end{array} \right]_{2 \times 3}$	3	lengths	
similarity	$\left[\begin{array}{c c} s\mathbf{R} & \mathbf{t} \end{array} \right]_{2 \times 3}$	4	angles	
affine	$\left[\begin{array}{c} \mathbf{A} \end{array} \right]_{2 \times 3}$	6	parallelism	
projective	$\left[\begin{array}{c} \tilde{\mathbf{H}} \end{array} \right]_{3 \times 3}$	8	straight lines	

(\mathbf{x} is in *homogeneous coordinates*)



Parametric (Global) Warping

Examples of parametric warps:

```
In [4]: im = cv2.imread('/Users/yiftachedelstain/ee046746-computer-vision/assets/tut_8_exm.jpg')
im = cv2.cvtColor(im, cv2.COLOR_BGR2RGB)
rows, cols, d = im.shape
print(im.shape)

# Translation:
tx,ty = [10,20]
h_T = np.float32([[1,0,tx],[0,1,ty],[0,0,1]])

# Rotation
theta = np.deg2rad(20)
h_R = np.float32([[np.cos(theta), -np.sin(theta), 0], [np.sin(theta), np.cos(theta), 0], [0,0,1]])

# Affine
h_AFF = np.array([[ 1.26666667,-0.5,-60], [-0.33333333,1,66.66666667],[0,0,1]])

H = [h_T,h_R,h_AFF]

img = [im]
for h in H:
    img.append(cv2.warpPerspective(im, h, (cols,rows)))
Titles = ['Original','Translation','Rotation','Affine']

# Perspective
pts1 = np.float32([[100,77],[320,105],[100,150],[385,170]])
pts2 = np.float32([[0,0],[300,0],[0,100],[300,50]])
M = cv2.getPerspectiveTransform(pts1,pts2)
im_perspective = cv2.warpPerspective(im,M, (400,300))

# Cylindrical (non linear)
def cylindricalWarp(img, K):
    """This function returns the cylindrical warp for a given image and intrinsics matrix
    h_,w_ = img.shape[:2]
    # pixel coordinates
    y_i, x_i = np.indices((h_,w_))
    X = np.stack([x_i,y_i,np.ones_like(x_i)],axis=-1).reshape(h_*w_,3) # to homog
    Kinv = np.linalg.inv(K)
    X = Kinv.dot(X.T).T # normalized coords
    # calculate cylindrical coords (sin\theta, h, cos\theta)
    A = np.stack([np.sin(X[:,0]),X[:,1],np.cos(X[:,0])],axis=-1).reshape(w_*h_,3)
    B = K.dot(A.T).T # project back to image-pixels plane
    # back from homog coords
    B = B[:, :-1] / B[:, [-1]]
    # make sure warp coords only within image bounds
    B[(B[:,0] < 0) | (B[:,0] >= w_) | (B[:,1] < 0) | (B[:,1] >= h_)] = -1
    B = B.reshape(h_,w_,-1)

    img_rgba = cv2.cvtColor(img, cv2.COLOR_RGB2RGBA) #BGR2BGRA for transparent borders...
    # warp the image according to cylindrical coords
    return cv2.remap(img_rgba, B[:, :, 0].astype(np.float32), B[:, :, 1].astype(np.float32),
                     cv2.INTER_AREA, borderMode=cv2.BORDER_TRANSPARENT)
# Camera parameters:
K = np.array([[200,0,cols/2],[0,400,rows/2],[0,0,1]]) # mock intrinsics
img_cyl = cylindricalWarp(im, K)

(362, 484, 3)
```

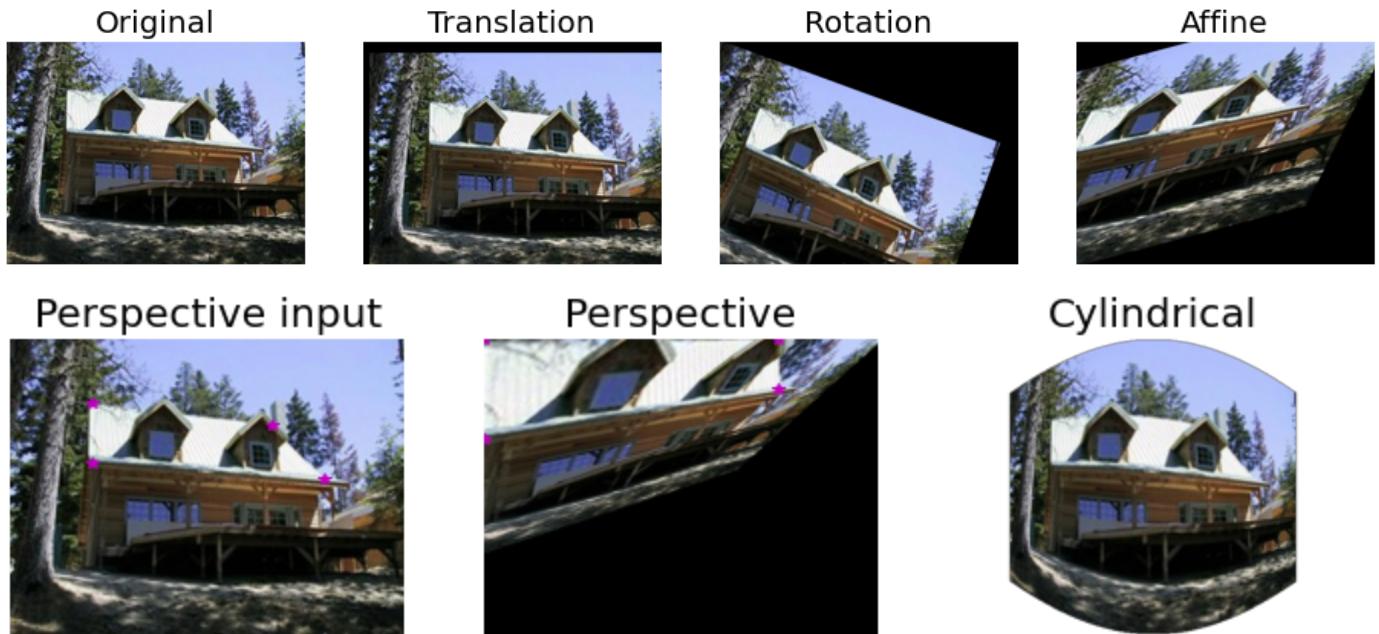
```
In [5]: plot_images(img,Titles,(1,4),figsize=(16,4),fontsize=20)

plt.figure(figsize=(12,4))
plt.subplot(131)
plt.imshow(im,plt.title('Perspective input',fontsize=20)
plt.plot(pts1[:,0],pts1[:,1],'m*')
plt.axis('off')
```

```

plt.subplot(132), plt.imshow(im_perspective), plt.title('Perspective', fontsize=20)
plt.plot(pts2[:,0], pts2[:,1], 'm*')
plt.axis('off')
plt.subplot(133), plt.imshow(img_cyl)
plt.title('Cylindrical', fontsize=20)
_ = plt.axis('off')

```



- [Code source](#) - OpenCV
- [Cylinder code source](#) - More Than Technical
- [Cylinder coordinates](#)

Basic 2D Transformations

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos\Theta & -\sin\Theta & 0 \\ \sin\Theta & \cos\Theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

Scale

Rotate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad \begin{bmatrix} x' \\ y' \\ w \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Affine

Projective
(Homography)

When can we use a homography matrix to describe the transformation of the scene?

Homography maps between:

- Points that lay on a plane in the real world and their positions in an image.
- Points in two different images that lay on the **same** plane.
- Two images of a 3D object where the camera has been rotated but has not been translated.

Homographies are used to approximate the transformations of far away objects:

- Works fine for small viewpoint changes.



Excercise 1 (Spring 2020 Q1A)

Write down the homography H (3×3), that maps from image-1 to image-2 in the following cases:

1. Image-2 is a translated image 1 by 100 pixels to the left.
2. Image-2 is a $3 \times$ zoom of image 1, so that the pixel indexed $(0, 0)$ stays in its original position.

Answer:

1. Image-2 is a translated image 1 by 100 pixels to the left.

We know that for a translation the homography matrix becomes:

$$\begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix}$$

The translation of the image is to the left, which means

$$x_2 = x_1 - 100$$

and the position in the y axis stays the same:

$$y_2 = y_1$$

Therefore, the homography that describes the given translation is given by:

$$\begin{bmatrix} 1 & 0 & -100 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Sanity check:

$$\begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & -100 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix} = \begin{bmatrix} x_1 - 100 \\ y_1 \\ 1 \end{bmatrix}$$

```
In [6]: im = cv2.imread('/Users/yiftachedelstain/ee046746-computer-vision/assets/tut_8_exm.jpg')
im = cv2.cvtColor(im, cv2.COLOR_BGR2RGB)
rows, cols, d = im.shape
h_T = np.float32([[1, 0, -100],
                  [0, 1, 0],
```

```

[0, 0, 1]])
im_T = cv2.warpPerspective(im, h_T, (cols,rows))
Titles = ['Original','Translation']

plot_images([im,im_T],Titles,(1,2),figsize=(16,4),fontsize=20)

```

Original



Translation



1. Image-2 is a $3 \times$ zoom of image 1, so that the pixel indexed $(0, 0)$ stays in its original position.

Answer:

The scaled image is different by 1 DOF, since we scale the X-axis and the Y-axis the same amount (same zoom). Going back to the homography matrix, this means only the s parameter changes:

$$\begin{bmatrix} 3 & 0 & 0 \\ 0 & 3 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

In [7]:

```

h_S = np.float32([[3,0,0],[0,3,0],[0,0,1]])
im_S = cv2.warpPerspective(im, h_S, (cols,rows))
Titles = ['Original','ZOOM']

plot_images([im,im_S],Titles,(1,2),figsize=(16,4),fontsize=20)

```

Original



ZOOM



Computing Parametric Transformations

- [Affine](#)
- [Projective](#)

Computing Affine Transformation

- Assuming we know correspondences, how do we get the transformation?

$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$
$$\begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix} = \begin{bmatrix} \dots & \dots & \dots & 0 & 0 & 0 \\ x_i & y_i & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_i & y_i & 1 \\ \dots & \dots & \dots & \dots & \dots & \dots \end{bmatrix} \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \\ h_5 \\ h_6 \end{bmatrix}$$
$$b = Ah$$

- Solve with Least-squares $\|Ah - b\|^2$

$$h = (A^T A)^{-1} A^T b$$

In Python:

```
A_inv = pinv(A)  
  
h = np.linalg.pinv(A)@b
```

- How many matches (correspondence pairs) do we need to solve?
- Once we have solved for the parameters, how do we compute the coordinates of the corresponding point for any pixel (x_{new}, y_{new}) ?

$$H = \begin{bmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ 0 & 0 & 1 \end{bmatrix}$$

$$x'_{new} = Hx_{new}$$

Computing Projective Transformation

- Recall working with homogenous coordinates

$$\begin{bmatrix} u_i \\ v_i \\ w_i \end{bmatrix} = \begin{bmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ h_7 & h_8 & h_9 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

$$x'_i = u_i/w_i$$

$$y'_i = v_i/w_i$$

- We get the following non-linear equation:

$$x'_i = \frac{h_1 x_i + h_2 y_i + h_3}{h_7 x_i + h_8 y_i + h_9}$$

$$y'_i = \frac{h_4 x_i + h_5 y_i + h_6}{h_7 x_i + h_8 y_i + h_9}$$

- We can re-arrange the equation

$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & 0 & \dots & -x_i x'_i & -y_i x'_i & -x'_i \\ 0 & 0 & 0 & x_i & y_i & 1 & \dots & -x_i y'_i & -y_i y'_i & -y'_i \end{bmatrix} \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \\ h_5 \\ h_6 \\ h_7 \\ h_8 \\ h_9 \end{bmatrix} = \begin{bmatrix} \dots \\ 0 \\ 0 \\ \dots \end{bmatrix}$$

- We want to find a vector h satisfying

$$Ah = 0$$

where A is full rank. We are obviously not interested in the trivial solution $h = 0$ hence we add the constraint

$$\|h\| = 1$$

- Thus, we get the homogeneous Least square equation:

$$\arg \min_h \|Ah\|_2^2, \text{ s.t. } \|h\|_2^2 = 1$$

Compute Projective transformation using SVD:

$$\arg \min_h \|Ah\|_2^2, \text{ s.t. } \|h\|_2^2 = 1$$

- Let decompose A using SVD: $A = UDV^T$, where U and V are orthonormal matrix, and D is a diagonal matrix.
 - Need a reminder on SVD? [Click Here](#)
- From orthonormality of U and V :

$$\|UDV^T h\| = \|DV^T h\|$$

$$\|V^T h\| = \|h\|$$

Hence, we get the following minimization problem:

$$\arg \min_h \|DV^T h\| \text{ s.t. } \|V^T h\| = 1$$

- Substitute $y = V^T h$:

$$\arg \min_h ||Dy|| \text{ s.t. } ||y|| = 1$$

- D is a diagonal matrix with decreasing values. Then, it is clear that $y = [0, 0, \dots, 1]^T$.

- Therefore, choosing h to be the last column in V will minimize the equation.

In Python:

```
(U,D,Vh) = np.linalg.svd(A, False)
```

```
h = Vh.T[:, -1]
```

Some more options to find h :

- Lagrange multipliers - [Least-squares Solution of Homogeneous Equations](#)
- Using EVD (eigenvalue decomposition) on $A^T A$.
- If we know our transformation is nearly Affine we can get an approximate solution using linear least squares
- As we said, some of the points may be a good match, and some won't.
- If we use a wrong point in the algorithm, we will get the wrong transformation.
- How can we solve that major problem?
- We can use our prior knowledge on the type of the trasformation, and use points that have strong evidence that are correct.
- For example, we know that homographys preserve lines, so we can try to look for keypoints that lay in a line, and use them to find the transformation.
- How would we find keypoints that lay on a line?



RANSAC

RANdom SAMple Consensus [Fischler & Bolles 1981]

-
- Key ideas:
 - Look for “inliers” and use only them
 - If we fit a model to “outliers” we will not get a good fitting

RANSAC Algorithm

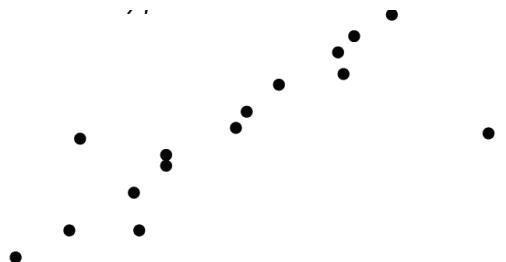
Loop:

1. Randomly select a group of points
2. Fit a model to the selected group
3. Find the inliers of the computed model
4. If number of inliers is large enough re-compute model using only inliers
5. Compute number of inliers of updated model

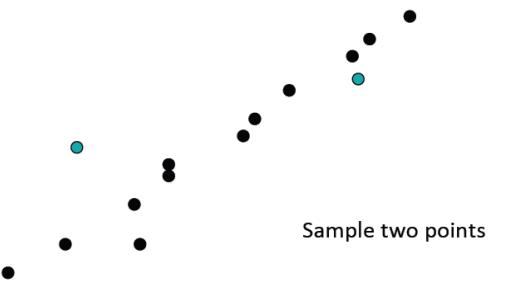
The winner: model with the largest number of inliers

Line Fitting with RANSAC

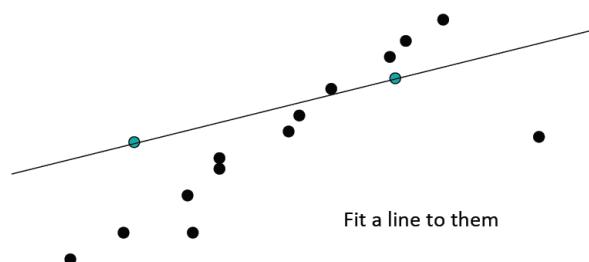
- Input: A set of edge points



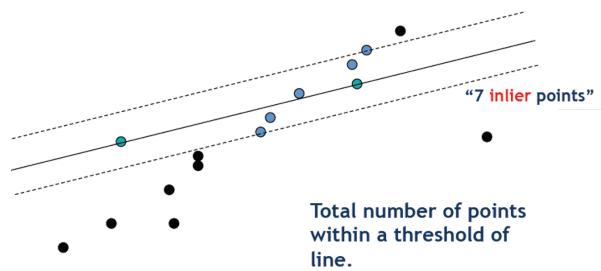
- **Step 1:** Select two points



- **Step 2:** Fit a line to the selected points



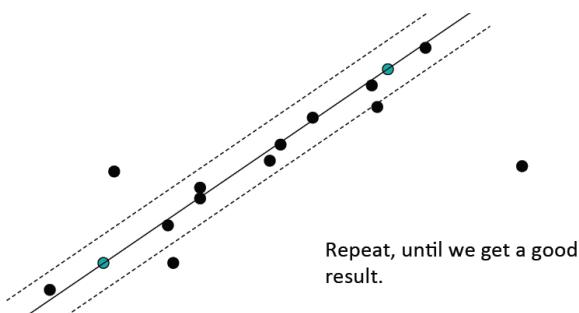
- **Step 3:** Identify inliers



- **Step 4:** Repeat Steps 1-3 until convergence

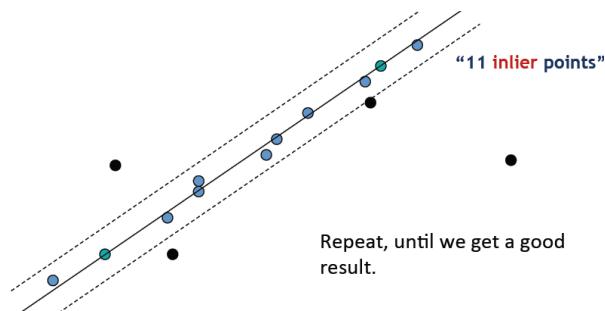
- **Iteration 2:** Steps 1-2:

- Sample new pts and fit a line



- **Iteration 2:** Step 3:

- Count number of new inliers



RANSAC – Stopping Criteria

- **Option 1:** when the model is good enough:

- By the number of inliers
- By its fitting error

- **Option 2:** according to probability

- Let K be the number of iterations
- Let n be the number of points needed to compute the model
- Let f be the fraction of inliers of a model
- Then the probability that a single sample is correct: f^n

- The probability that all K samples fail is $(1 - f^n)^K$
- Choose K high enough to keep the failure rate low enough

RANSAC – For Multiple Models?

- How can we use RANSAC to compute multiple models?
 - Apply RANSAC multiple times, each time we remove the detected inliers.

RANSAC - Summary

- **Pros**

- General method that works well for lots of model fitting problems
- Easy to implement

- **Cons**

- When the percentage of outliers is high too many iterations are needed and failure rate increases

- **Implementation**

- OpenCV usage with Homographies (Tutorial 7): `cv2.findHomography(src_pts, dst_pts, cv2.RANSAC)`
- Scikit-image: `skimage.measure.ransac()`
- Scikit-learn: `sklearn.linear_model.RANSACRegressor()`
- and many more..

Many improved algorithms:

- **PROSAC**
 - Key idea is to assume that the similarity measure predicts correctness of a match
- **Randomized RANSAC**
 - Each step take a random subset of the query points and perform RANSAC
- **KALMANSAC**
- **and more...**
- Estimating homogrphy with RANSAC in OpenCV: `cv2.findHomography(src_pts, dst_pts, cv2.RANSAC)`



Excercise 2 (Spring 2020 Q1D) - Ransac

We'd like to fit an ellipse to a given set of 200 points, $X_i = (x_i, y_i), i = 1, \dots, 200$.

As you know, an ellipse can be described by the implicit equation $a^T(x^2, xy, y^2, x, y, 1) = 0$ where $a = (a_1, a_2, a_3, a_4, a_5, a_6)$ is a coefficients vector. The data points might be noisy, such that some of the points may not reside on the perimeter of the ellipse (outliers).

- Describe the RANSAC algorithm needed to fit the ellipse (and find a).

Answer:

- to define an ellipse, we need at least 5 points ([Why?](#)). This means a is defined up to a scale factor.

Step 1: Sample 5 points from X_i randomly. **Step 2:** Fit an ellipse to the points with Least Squares.

•

$$\begin{bmatrix} & & \dots & & & \\ x_i^2 & x_i y_i & y_i^2 & x_i & y_i & 1 \\ & & \dots & & & \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \\ a_5 \\ a_6 \end{bmatrix} = 0$$

- This can be denoted as $Sa = 0$, and solved via SVD under the constraint $\|a\| = 1$

Step 3: Find the points that fit to the model (inliers), i.e. points with a fit error below a threshold chosen in advanced.

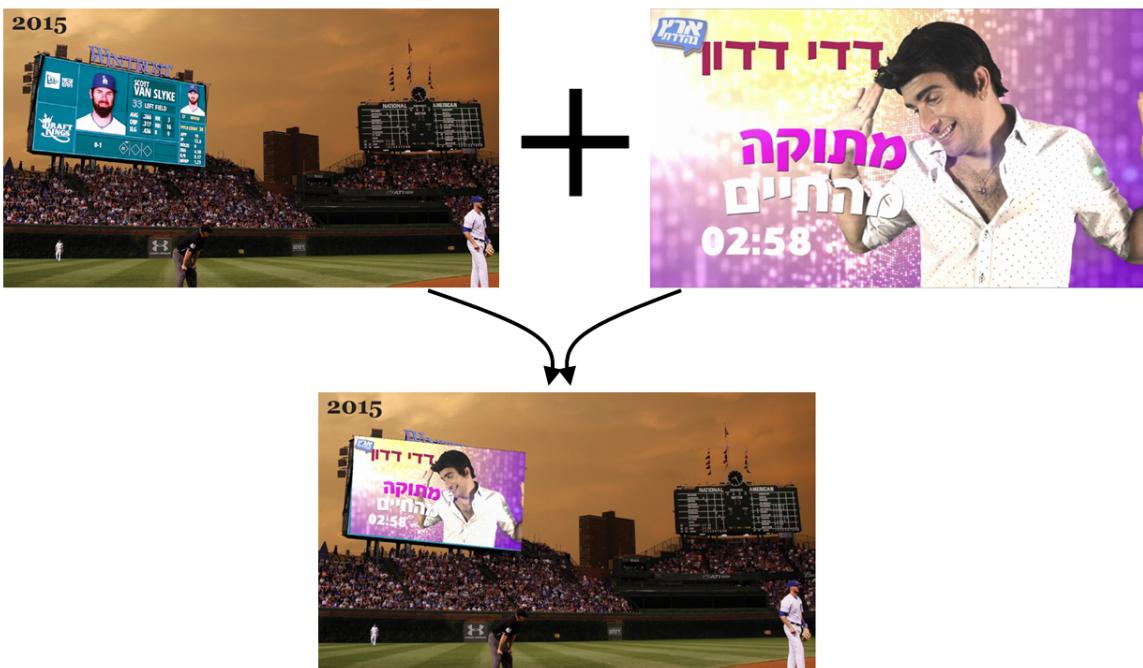
Step 4: Refine our model according to all the inliers, and count them. The number of inliers is the "score" of the current model (from this iteration).

Step 5: Repeat 1-4, until the stop criteria has been met.

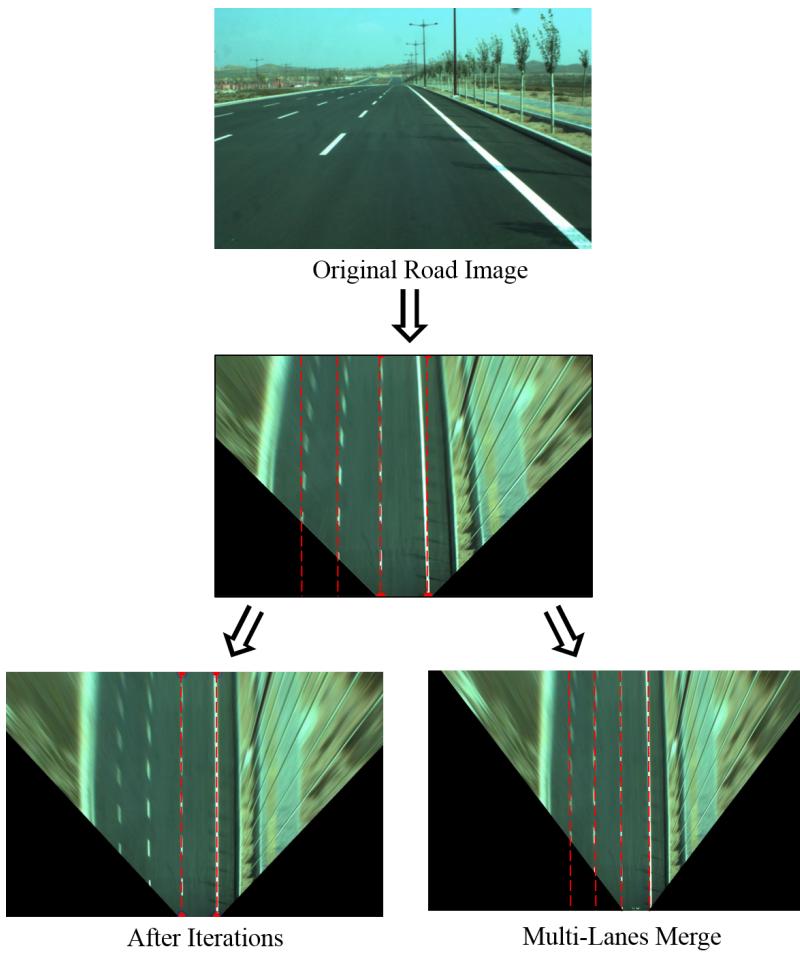
When do we stop?

- The percentage of inliers out of all the points is above a certain percentage.
- The amount of inliers is above a certain amount.
- An error threshold (for the LS) is reached.
- A predefined amount of iterations is reached.

application 1: Planting images into other images

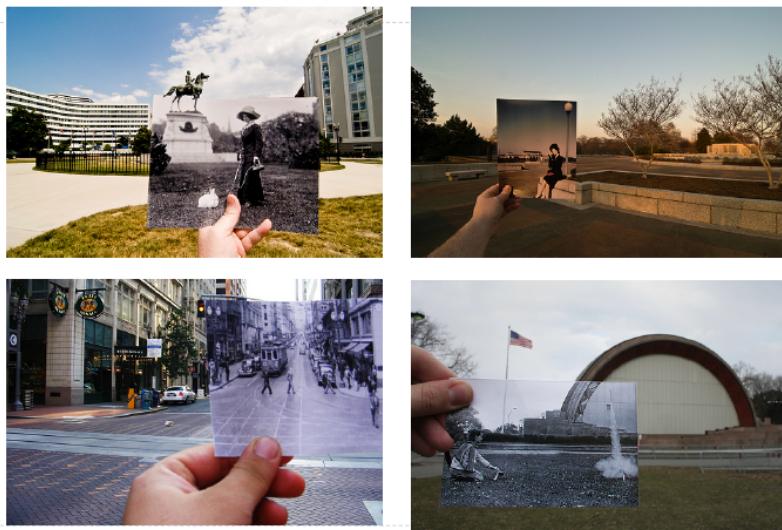


application 2: BirdEye



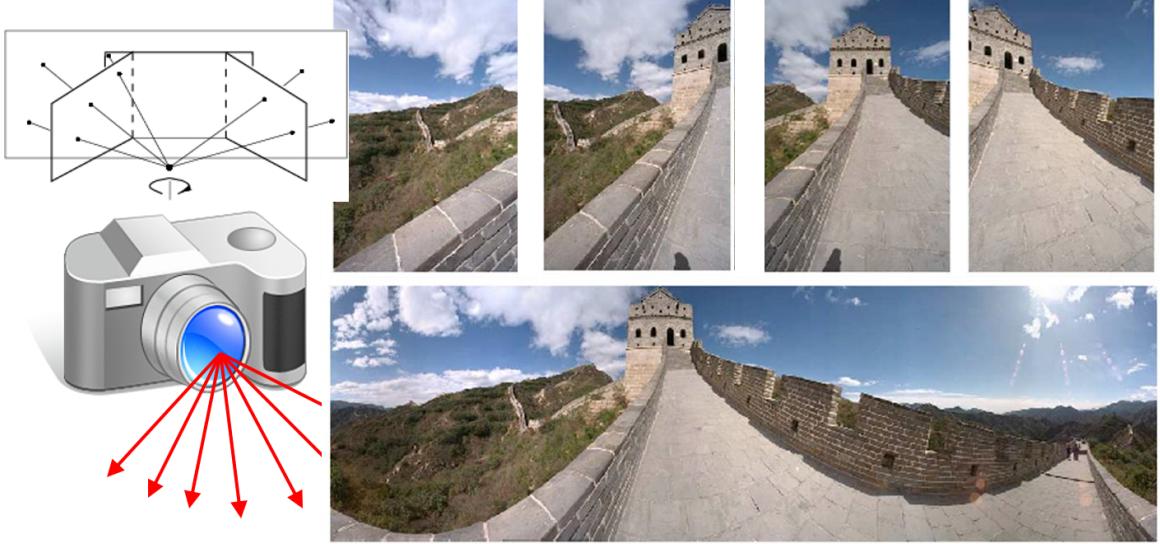
- [Image Source](#)

application 3: Looking into the past



- [Image Source from Flickr](#)

application 4: Panorama

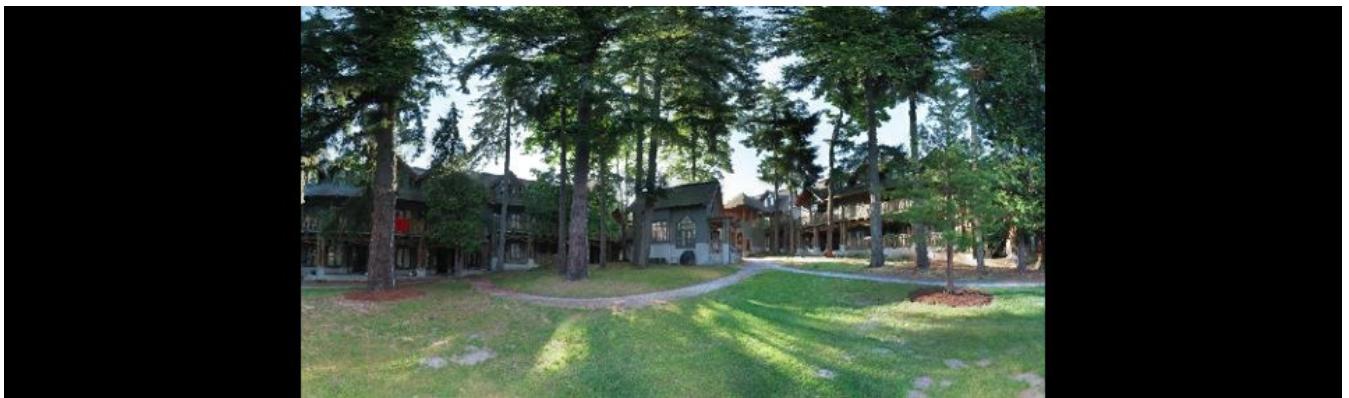
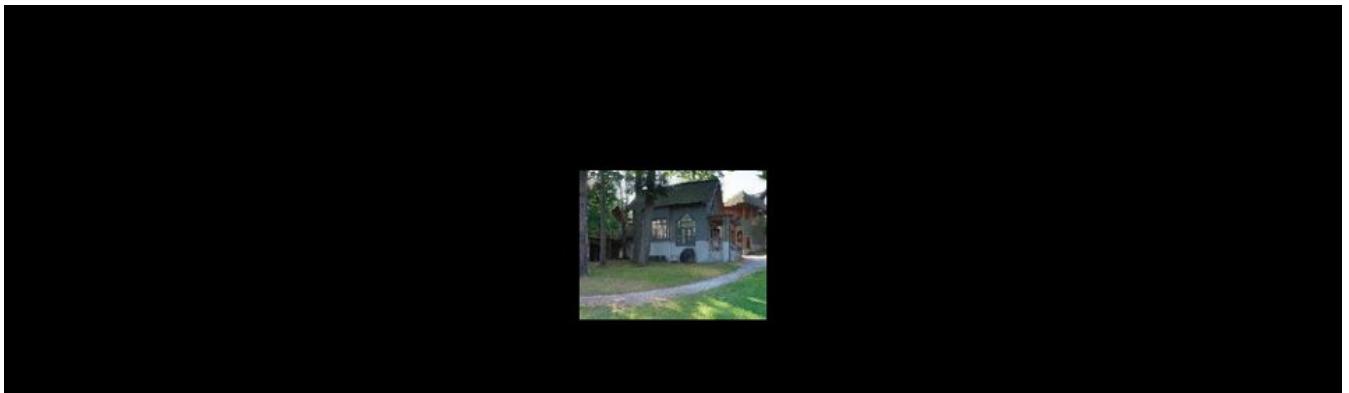


Obtain a wider angle view by combining multiple images



Panorama

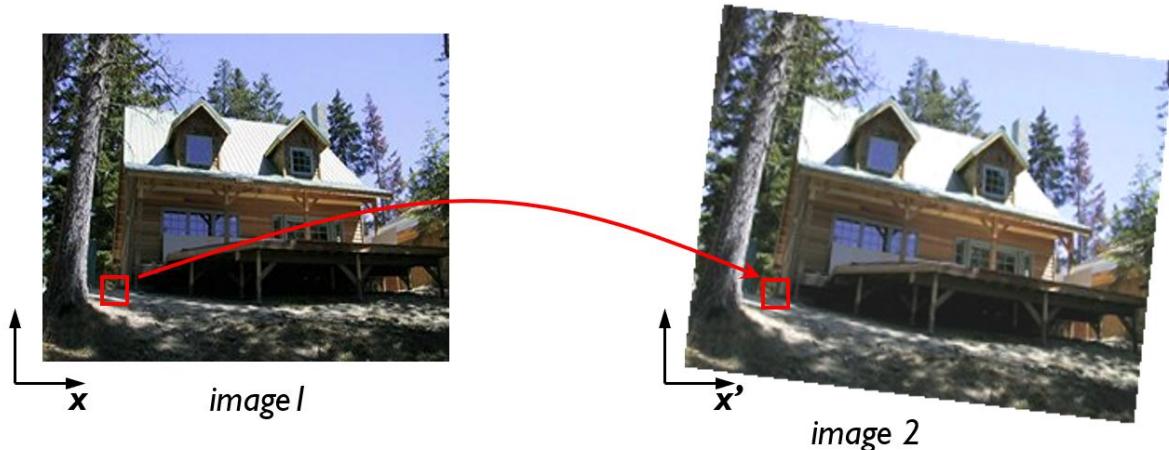
- Warping
- Image Blending (Feathering)





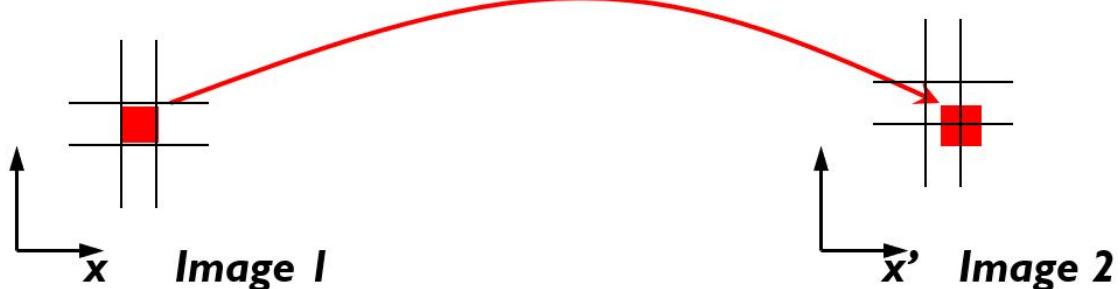
Warp - What we need to solve?

- Given source and target images, and the transformation between them, how do we align them?
- Send each pixel x in image1 to its corresponding location x' in image 2

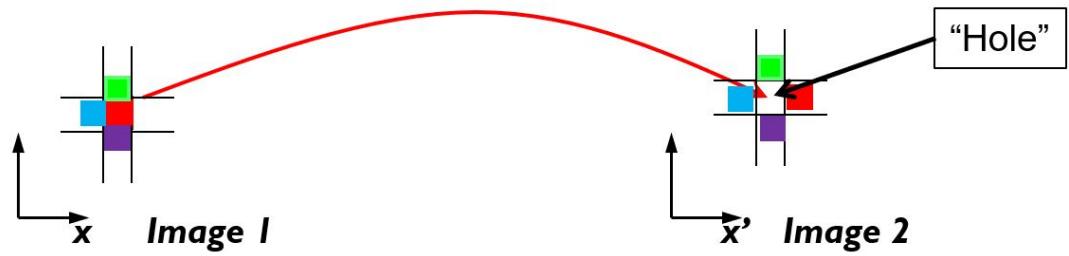


Forward Warping

- What if pixel lands “between” two pixels?
- Answer: add “contribution” to several pixels and normalize (splatting)

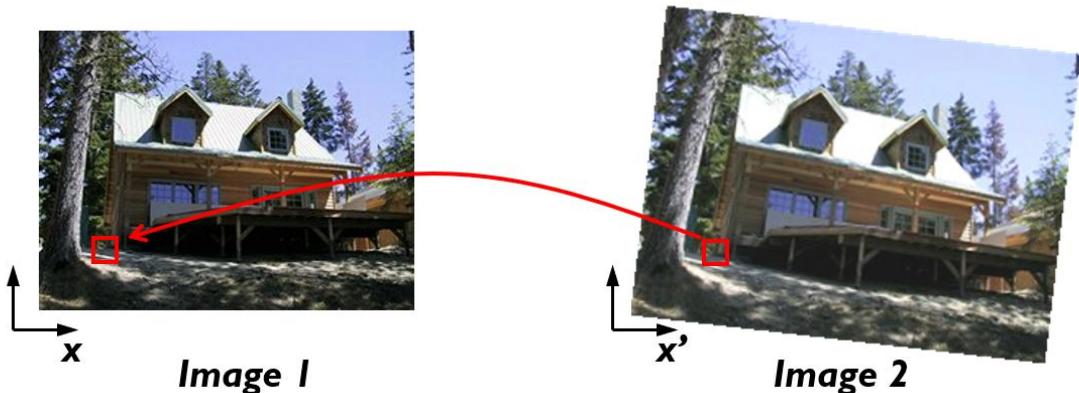


- Limitation: Holes (some pixels are never visited)



Inverse Warping

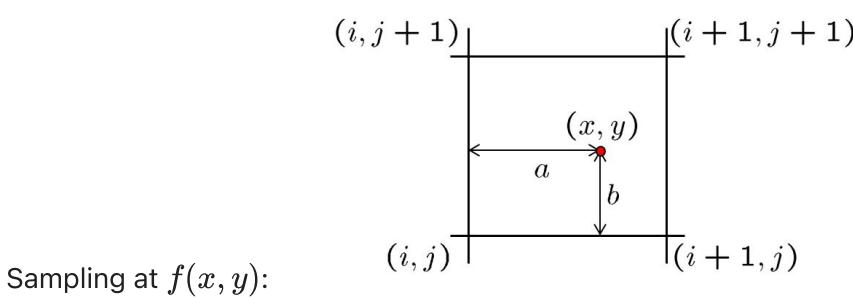
- For each pixel x' in image 2 find its origin x in image 1
- Problem: What if pixel comes from "between" two pixels?



- Answer: interpolate color value from neighbors



Bilinear Interpolation



Sampling at $f(x, y)$:

$$\begin{aligned} f(x, y) = & (1 - a)(1 - b)f[i, j] \\ & + a(1 - b)f[i + 1, j] \\ & + abf[i + 1, j + 1] \\ & + (1 - a)bf[i, j + 1] \end{aligned}$$

Python:

- `interp2d()` - <https://docs.scipy.org/doc/scipy/reference/generated/scipy.interpolate.interp2d.html>
- Inverse warping in OpenCV: `cv2.warpPerspective(im,*,*,cv2.WARP_INVERSE_MAP)`



Excercise 3 (Spring 2020 Q1B)

1. Describe forward and backward warping.
2. For which out of the two cases before will there be a difference between using inverse or forward warping?
3. For the case which a difference exists, what type of warping should be used and why?

reminder:

- Image-2 is a translated image 1 by 100 pixels to the left
- Image-2 is a 3x zoom of image 1, so the index (0,0) stay in place

Answer

1. Read the tutorial :)
2. For the translation there will be no difference between forward/inverse warping, while for the zooming there will be a difference - we'll get holes.

To understand why, let's denote the pixel of image k at the indices (i, j) as x_{ij}^k .

In the first case, each pixel is moved to the left by a whole number, therefore for forward warping we follow the warp defined by:

$$x_{(i-100)j}^2 = x_{ij}^1$$

And for inverse warping we use:

$$x_{ij}^2 = x_{(i+100)j}^1$$

There's no need to interpolate anything here - because the pixel moved by an integer, and thus there is no difference.

In the second case we get for forward warping:

$$x_{(3i)(3j)}^2 = x_{ij}^1$$

And for inverse warping we get an interpolation dependant on α, β :

$$\begin{aligned} x_{ij}^2 &= (1 - \alpha)(1 - \beta)x_{\left\lfloor \frac{i}{3} \right\rfloor \left\lfloor \frac{j}{3} \right\rfloor}^1 + \alpha(1 - \beta)x_{\left(\left\lfloor \frac{i}{3} \right\rfloor + 1\right) \left\lfloor \frac{j}{3} \right\rfloor}^1 + \alpha\beta x_{\left\lfloor \frac{i}{3} \right\rfloor \left(\left\lfloor \frac{j}{3} \right\rfloor + 1\right)}^1 \\ &\quad + (1 - \alpha)\beta x_{\left(\left\lfloor \frac{i}{3} \right\rfloor + 1\right) \left(\left\lfloor \frac{j}{3} \right\rfloor + 1\right)}^1 \end{aligned}$$

We can see that for the forward case we get "holes" in image 2, for example $x_{1,1}^2$ has no origin in image 1.

1. The backward wrapping will be of course better, since as we saw no holes will be created.

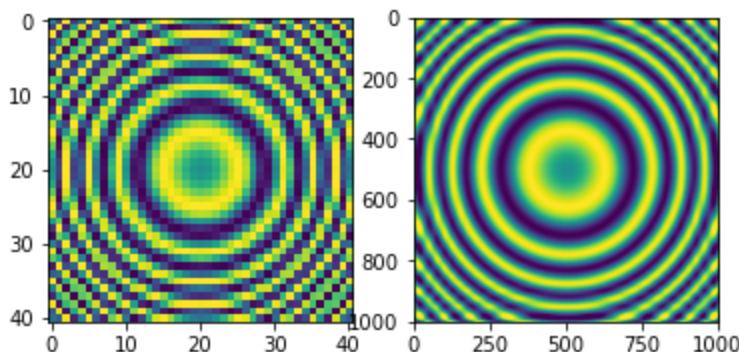
```
In [8]: import numpy as np
from scipy import interpolate
import matplotlib.pyplot as plt

# original samples
x = np.arange(-5.01, 5.01, 0.25)
y = np.arange(-5.01, 5.01, 0.25)
xx, yy = np.meshgrid(x, y)
z = np.sin(xx**2+yy**2)

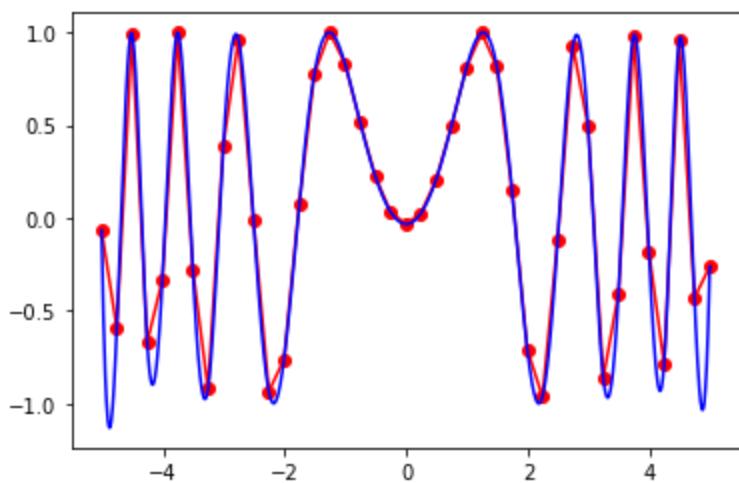
# interpolated samples
xnew = np.arange(-5.01, 5.01, 1e-2)
ynew = np.arange(-5.01, 5.01, 1e-2)
```

```
In [9]: ## Cubic
f = interpolate.interp2d(x, y, z, kind='cubic')
znew = f(xnew, ynew)

plt.figure()
plt.subplot(121)
plt.imshow(z, vmin=-1, vmax=1)
plt.subplot(122)
plt.imshow(znew, vmin=-1, vmax=1)
plt.show()
```

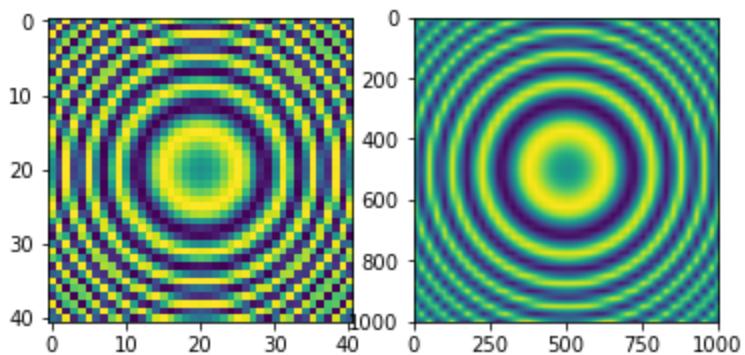


```
In [10]: plt.figure()
plt.plot(x, z[0, :], 'ro-', xnew, znew[0, :], 'b-')
plt.show()
```

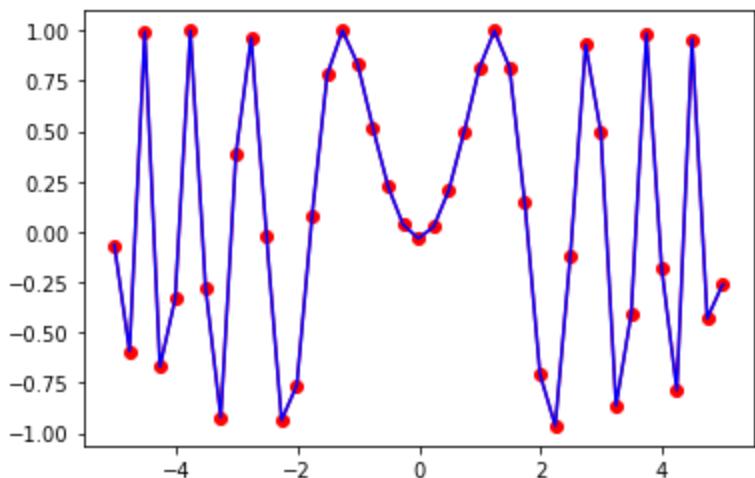


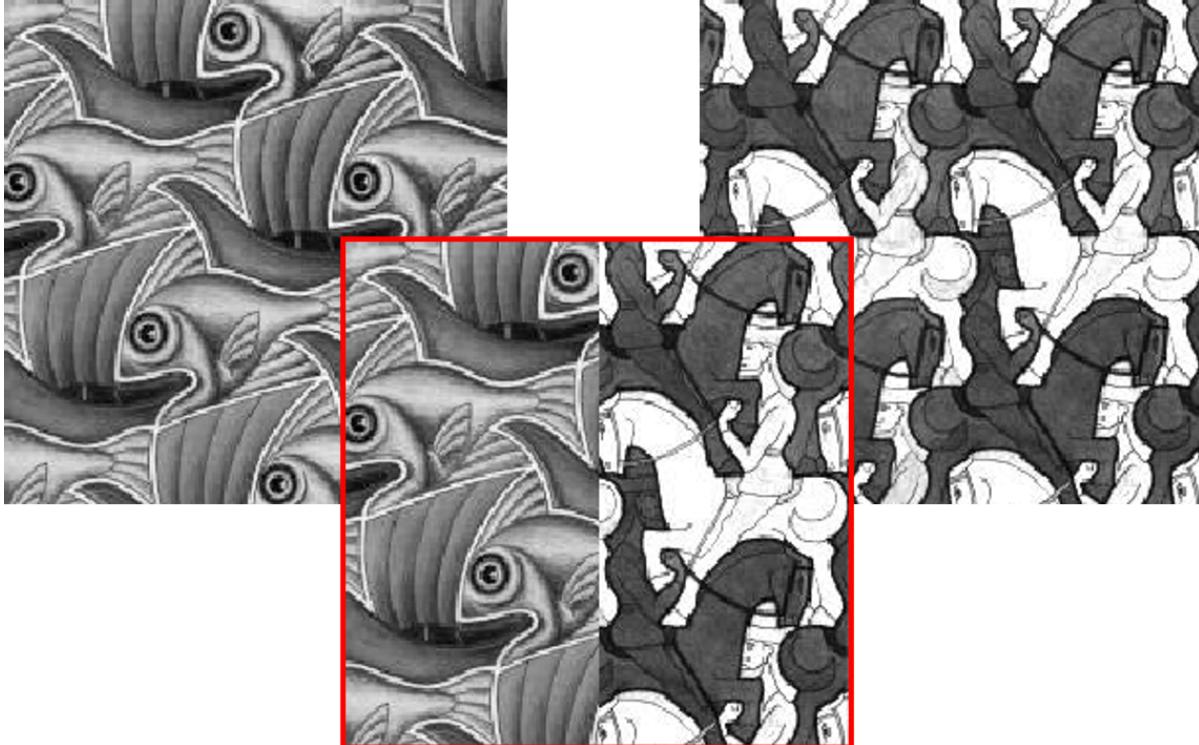
```
In [11]: ##linear
f2 = interpolate.interp2d(x, y, z, kind='linear')
znew2 = f2(xnew, ynew)

plt.figure()
plt.subplot(121)
plt.imshow(z, vmin=-1, vmax=1)
plt.subplot(122)
plt.imshow(znew2, vmin=-1, vmax=1)
plt.show()
```



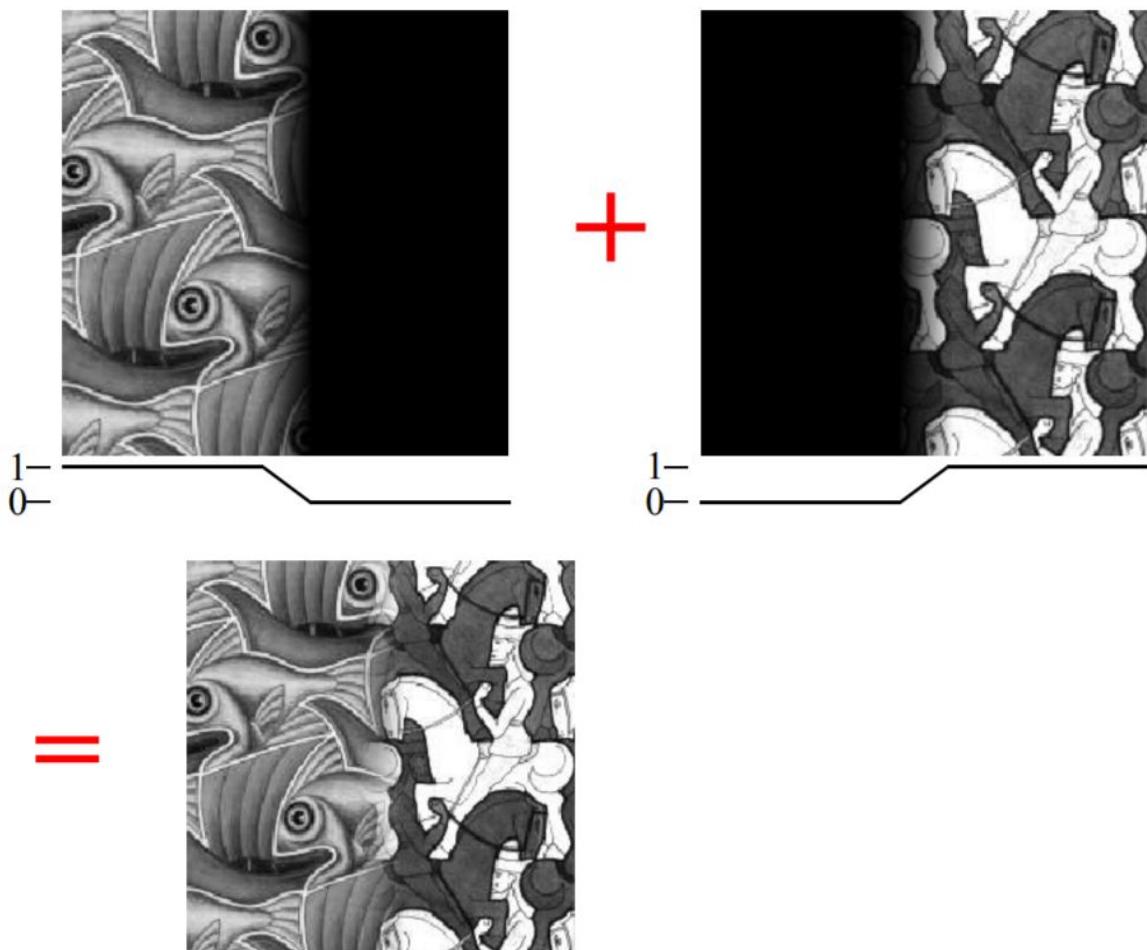
```
In [12]: plt.figure()
plt.plot(x, z[0, :], 'ro-', xnew, znew2[0, :], 'b-')
plt.show()
```





- Alpha blending
- Pyramid blending

Alpha Blending



Pyramid Blending:

1. Build a Gaussian pyramid for each image
2. Build the Laplacian pyramid for each image
3. Decide/find the blending border (in the example: left half belongs to image 1, and right half to image 2 -> the blending border is `cols/2`)
 - Split by index, or
 - Split using a 2 masks (can be weighted masks)
4. Construct a new mixed pyramid - mix each level separately according to (3)
5. Reconstruct a blend image from the mixed pyramid

```
In [13]: def create_pyrs (A,B) :
    # generate Gaussian pyramid for A
    G = A.copy()
    gpA = [G]
    for i in range(6):
        G = cv2.pyrDown(G)
        gpA.append(G)
    # generate Gaussian pyramid for B
    G = B.copy()
    gpB = [G]
    for i in range(6):
        G = cv2.pyrDown(G)
        gpB.append(G)
    # generate Laplacian Pyramid for A
    lpA = [gpA[5]]
    for i in range(5,0,-1):
        GE = cv2.pyrUp(gpA[i])
        L = cv2.subtract(gpA[i-1],GE)
        lpA.append(L)
    # generate Laplacian Pyramid for B
    lpB = [gpB[5]]
    for i in range(5,0,-1):
        GE = cv2.pyrUp(gpB[i])
        L = cv2.subtract(gpB[i-1],GE)
        lpB.append(L)
    return lpA,lpB
```

```
In [1]: def blend_images (A,B) :
    lpA,lpB = create_pyrs (A,B)
    # Now add left and right halves of images in each level
    LS = []
    for la,lb in zip(lpA,lpB):
        rows,cols,dpt = la.shape
        ls = np.hstack((la[:,0:int(cols/2)], lb[:,int(cols/2):])) #mixing can also be done
        LS.append(ls)
    # now reconstruct
    ls_ = LS[0]
    for i in range(1,6):
        ls_ = cv2.pyrUp(ls_)
        ls_ = cv2.add(ls_, LS[i])
    # image with direct connecting each half
    real = np.hstack((A[:,0:int(cols/2)],B[:,int(cols/2):]))
    return real, ls_
```

```
In [15]: def switch_texture (A,B) :
    lpA,lpB = create_pyrs (A,B)
    # Now add left and right halves of images in each level
    LS = []
    #     for la,lb in zip(lpA,lpB):
    #         rows,cols,dpt = la.shape
    #         ls = np.hstack((la[:,0:int(cols/2)], lb[:,int(cols/2):])) #mixing can also be done
```

```

#           LS.append(ls)
# now reconstruct
ls_ = lpA[0]
for i in range(1, 6):
    ls_ = cv2.pyrUp(ls_)
    ls_ = cv2.add(ls_, lpB[i])
# image with direct connecting each half
real = np.hstack((A[:, :int(cols/2)], B[:, int(cols/2):]))
return real, ls_

```

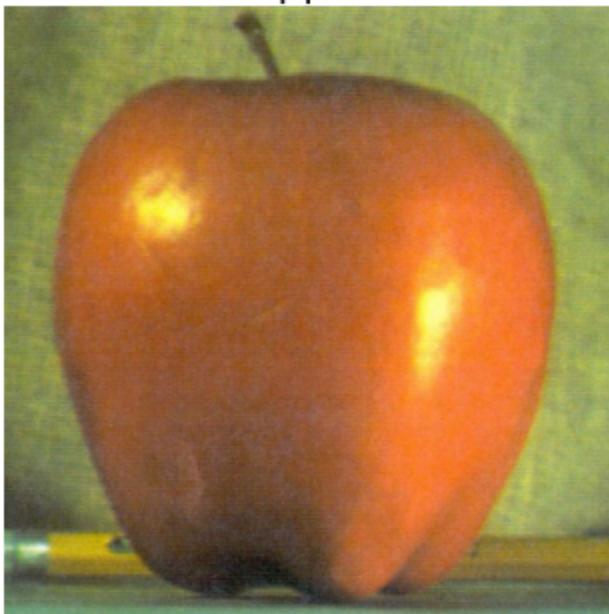
```
In [16]: def alpha_blend(A,B,MASK):
    A = np.float32(A)
    B = np.float32(B)
    return np.uint8(A*MASK), np.uint8(A*MASK+B*(1-MASK))
```

```
In [17]: A = cv2.imread('/Users/yiftachedelstain/ee046746-computer-vision/assets/apple.jpg')
B = cv2.imread('/Users/yiftachedelstain/ee046746-computer-vision/assets/orange.jpg')
real,ls_ = blend_images(A,B)

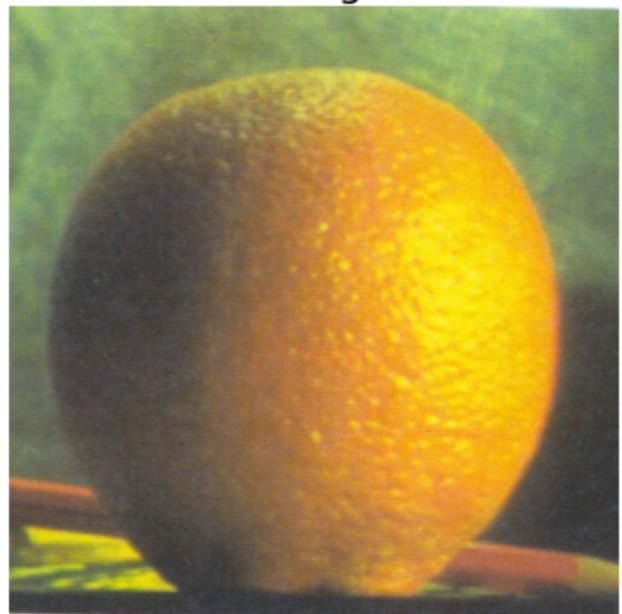
A_ = cv2.cvtColor(A,cv2.COLOR_BGR2RGB)
B_ = cv2.cvtColor(B,cv2.COLOR_BGR2RGB)
ls_ = cv2.cvtColor(ls_,cv2.COLOR_BGR2RGB)
real = cv2.cvtColor(real,cv2.COLOR_BGR2RGB)
```

```
In [18]: plot_images([A_,B_],['Apple','Orange'],(1,2),figsize=(12,12),fontsize=20)
```

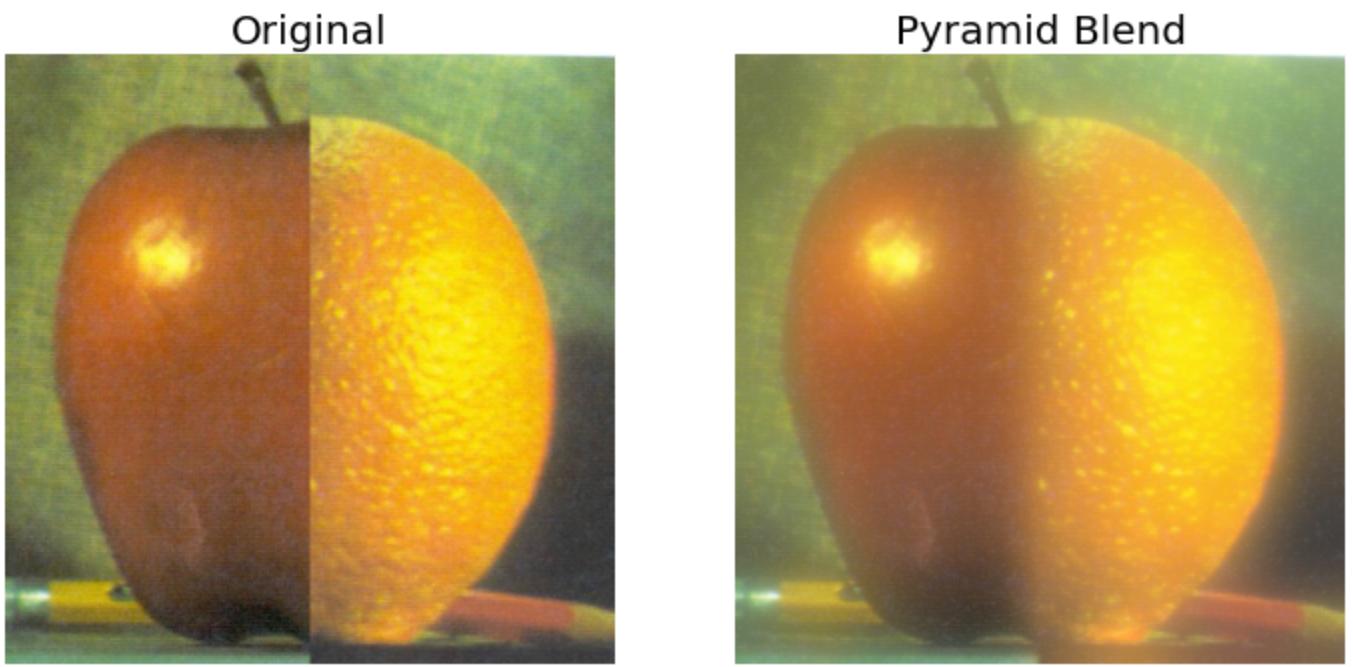
Apple



Orange



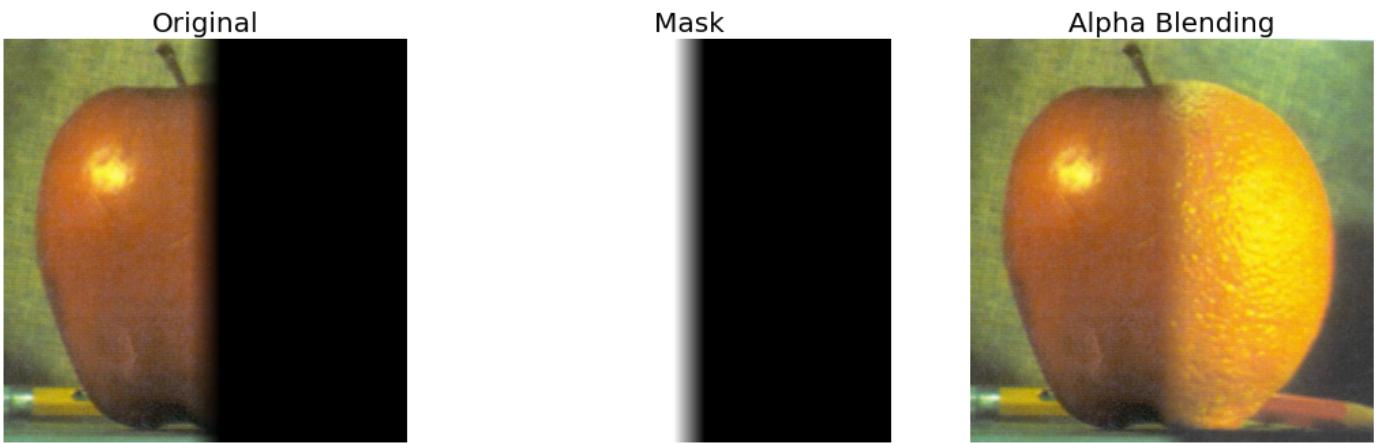
```
In [19]: plot_images([real,ls_],['Original','Pyramid Blend'],(1,2),figsize=(12,12),fontsize=20)
```



- Alpha blending example:

```
In [20]: MASK = np.ones_like(A,np.float32)
rows,cols,dpt = MASK.shape
w=20
v_dec = np.linspace(1,0,2*w)
MASK[:,int(cols/2):]=0
MASK[:,(int(cols/2)-w):(int(cols/2)+w)]=np.tile(np.reshape(v_dec,[1,-1,1]),[rows,1,3])
real,ls_ = alpha_blend(A,B,MASK)
ls_ = cv2.cvtColor(ls_,cv2.COLOR_BGR2RGB)
real = cv2.cvtColor(real,cv2.COLOR_BGR2RGB)
```

```
In [21]: plot_images([real,MASK,ls_],['Original','Mask','Alpha Blending'],(1,3),figsize=(18,12),f
```

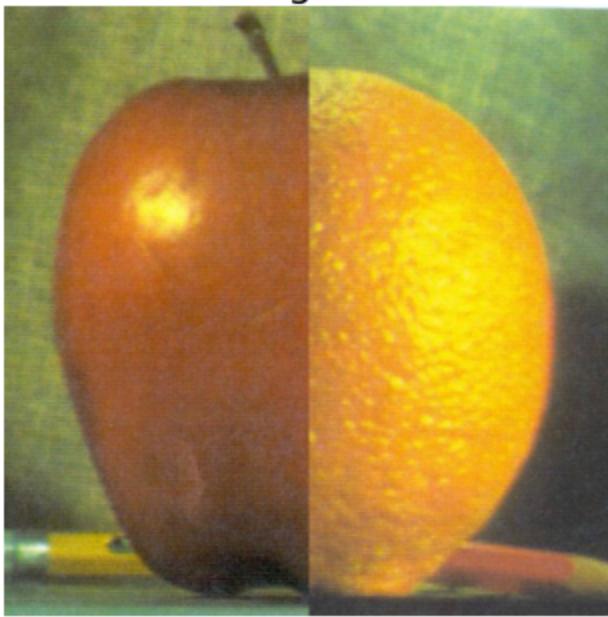


- Stylize image using pyramids:

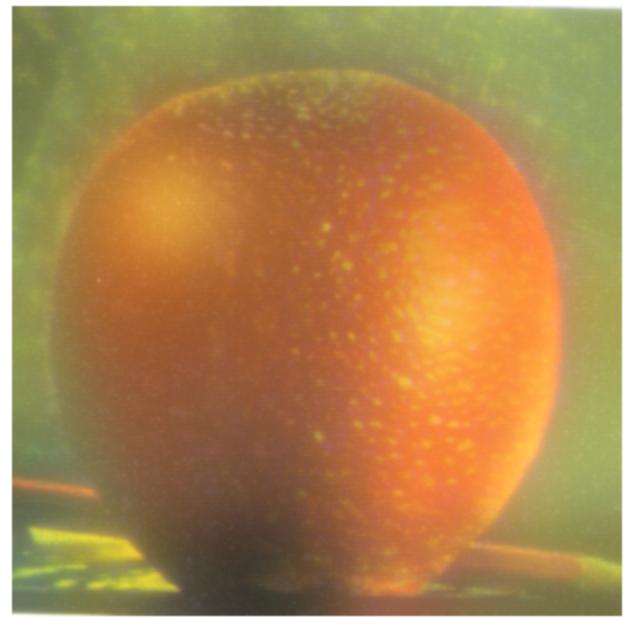
```
In [22]: real,ls_ = switch_texture(A,B)
ls_ = cv2.cvtColor(ls_,cv2.COLOR_BGR2RGB)
real = cv2.cvtColor(real,cv2.COLOR_BGR2RGB)
```

```
In [23]: plot_images([real,ls_],['Original','Texture'],(1,2),figsize=(12,12),fontsize=20)
```

Original



Texture



Blending Improvements

- Many algorithms have different variations of combining alpha and pyramid blending (different masks for different frequencies)
- Find the boundaries using **segmentation**



Panorama - Summary

- Detect features
- Compute transformations between pairs of frames
- Can Refine transformations using RANSAC
- Warp all images onto a single coordinate system
- Find mixing borders (e.g. using segmentation)
- Blend



Transformations in Deep Learning

- Can we incorporate transformations in the pipeline of a deep learning algorithm?
 - Moreover, can we accelerate these transformations by performing them on a GPU?
- **YES!**

Kornia - Computer Vision Library for PyTorch



- Kornia is a differentiable computer vision library for PyTorch
 - That means you can have gradients for the transformations!
- Inspired by OpenCV, this library is composed by a subset of packages containing operators that can be inserted within neural networks to train models to perform image transformations, epipolar geometry, depth estimation, and low-level image processing such as filtering and edge detection that operate directly on tensors.
- Check out `kornia.geometry` - <https://kornia.readthedocs.io/en/latest/geometry.html>
- Warp image using perspective transform

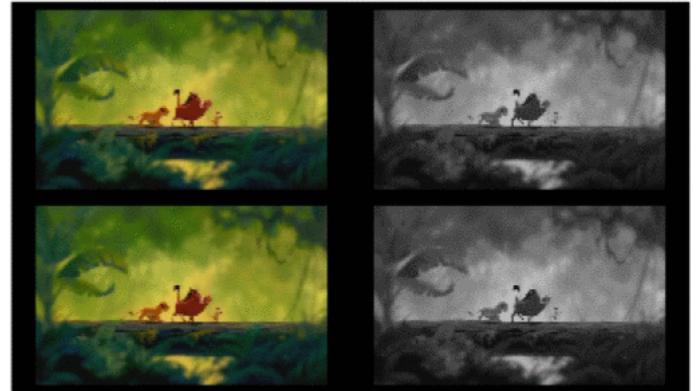


PyTorch

```
import torch
import kornia

frame: torch.Tensor = load_video_frame(...)

out: torch.Tensor = (
    kornia.rgb_to_grayscale(frame)
)
```



Recommended Videos



Warning!

- These videos do not replace the lectures and tutorials.
- Please use these to get a better understanding of the material, and not as an alternative to the written material.

Video By Subject

- Homography
 - Image geometry and planar homography - [ENB339 lecture 9: Image geometry and planar homography](#)
 - Homography - [Homography in computer vision explained](#)
- Transformations - [Lect. 5\(1\) - Linear and affine transformations](#)
- Matching Local Features
 - SIFT - [CSCI 512 - Lecture 12-1 SIFT](#)



Credits

- EE 046746 Spring 2022 - [Moshe Kimhi, Hila Manor](#)
- EE 046746 Spring 2020 - [Dahlia Urbach](#)

- Slides - Elad Osherov (Technion), Simon Lucey (CMU)
- Multiple View Geometry in Computer Vision - Hartley and Zisserman - Section 2
- [Least-squares Solution of Homogeneous Equations](#) - Center for Machine Perception - Tomas Svoboda
- [Computer vision: models, learning and inference](#) , Simon J.D. Prince - Section 15.1
- [Computer Vision: Algorithms and Applications](#) - Richard Szeliski - Sections 2,4,6, 9 (Free for Technion students via remote library)
- Icons from [Icon8.com](#) - <https://icons8.com>